

## PROGRAM 1 : (FIND S)

```
[10]: import csv  
with open('MLlab.csv','r')as cfe:  
    reader=csv.reader(cfe,delimiter=',')  
    train=[]  
    for row in reader:  
        train.append(row)  
print(train)
```

```
[['Sunny', 'Warm', 'Normal', 'Strong', 'Warm', 'S  
'High', 'Strong', 'Warm', 'Change', 'No'], ['Sunny',
```

```
[16]: hypo=['0']*6  
row=4  
col=6  
for i in range(row):  
    if(train[i][6]=="Yes"):  
        for j in range (col):  
            if(hypo[j]=='0'):br/>                hypo[j]=train[i][j]  
            else:  
                if(hypo[j]!=train[i][j]):  
                    hypo[j]='?'
```

```
[17]: print(hypo)
```

```
['Sunny', 'Warm', '?', 'Strong', '?', '?']
```

```
[ ]:
```

## PROGRAM 2 : (LINEAR)

```
[1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
[2]: df=pd.read_csv('kc_house_data.csv')  
df
```

	<b>id</b>	<b>date</b>	<b>price</b>	<b>l</b>
<b>0</b>	7129300520	20141013T000000	221900.0	
<b>1</b>	6414100192	20141209T000000	538000.0	
<b>2</b>	5631500400	20150225T000000	180000.0	
<b>3</b>	2487200875	20141209T000000	604000.0	
<b>4</b>	1954400510	20150218T000000	510000.0	
...	...	...	...	...
<b>21608</b>	263000018	20140521T000000	360000.0	
<b>21609</b>	6600060120	20150223T000000	400000.0	
<b>21610</b>	1523300141	20140623T000000	402101.0	

```
[3]: df=df.drop(['id','date'],axis=1)
print(df)
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
0	221900.0	3	1.00	1180	5650	1.0	
1	538000.0	3	2.25	2570	7242	2.0	
2	180000.0	2	1.00	770	10000	1.0	
3	604000.0	4	3.00	1960	5000	1.0	
4	510000.0	3	2.00	1680	8080	1.0	
...	...	...	...	...	...	...	...
21608	360000.0	3	2.50	1530	1131	3.0	
21609	400000.0	4	2.50	2310	5813	2.0	
21610	402101.0	2	0.75	1020	1350	2.0	
21611	400000.0	3	2.50	1600	2388	2.0	
21612	325000.0	2	0.75	1020	1076	2.0	
	waterfront	view	condition	grade	sqft_above	sqft_basement	\
0	0	0	3	7	1180	0	
1	0	0	3	7	2170	400	
2	0	0	3	6	770	0	
3	0	0	5	7	1050	910	
4	0	0	3	8	1680	0	
...	...	...	...	...	...	...	...
21608	0	0	3	8	1530	0	
21609	0	0	3	8	2310	0	
21610	0	0	3	7	1020	0	
21611	0	0	3	8	1600	0	
21612	0	0	3	7	1020	0	

```
[4]: y=df['price']
x=df.drop('price',axis=1)
x
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floo
0	3	1.00	1180	5650	1
1	3	2.25	2570	7242	2
2	2	1.00	770	10000	1
3	4	3.00	1960	5000	1
4	3	2.00	1680	8080	1
...	...	...	...	...	...
21608	3	2.50	1530	1131	3
21609	4	2.50	2310	5813	2
21610	2	0.75	1020	1350	2
21611	3	2.50	1600	2388	2
21612	2	0.75	1020	1076	2


21613 rows × 18 columns

```
[5]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=2)
```

```
[6]: from sklearn.linear_model import LinearRegression  
reg=LinearRegression().fit(x_train,y_train)  
reg.score(x_test,y_test)
```

```
[6]: 0.7156152938565365
```

```
[7]: #Predict values for test data  
y_pred = reg.predict(x_test)  
  
# Plot Actual vs Predicted  
plt.figure(figsize=(8,6))  
plt.scatter(y_test, y_pred)      # actual vs predicted points  
plt.plot(y_test, y_test)        # 45-degree line (perfect predictions)  
plt.xlabel("Actual House Price")  
plt.ylabel("Predicted House Price")  
plt.title("Actual vs Predicted House Prices (Linear Regression)")  
plt.show()
```



## PROGRAM 3 : (LOGISTIC)

```
[1]: import pandas as pd
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score,confusion_matrix
import numpy as np
import seaborn as sns
```

```
[2]: diabetes=load_diabetes()
x,y=diabetes.data,diabetes.target
y_binary=(y>np.median(y)).astype(int)
y_binary
```

```
[2]: array([1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0,
          0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0,
          1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1,
          1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0,
          0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0,
          0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1,
          0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1,
          1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1,
          1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0,
          1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0,
          1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1,
          0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1,
          0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0,
          1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0,
```

```
[3]: x_train,x_test,y_train,y_test=train_test_split(x,y_binary,test_size=0.2,random_state=42)
#print(len(x_test))
#print(len(x_train))
```

```
[4]: scaler=StandardScaler()
x_train=scaler.fit_transform(x_train)
x_test=scaler.fit_transform(x_test)
```

```
[5]: model=LogisticRegression()
model.fit(x_train,y_train)
x_train
```

```
[5]: array([[ 1.49836523,  1.06136988,  0.21990201, ...,  0.71103773,
       0.54748197, -0.06144896],
       [-0.22885822,  1.06136988, -0.41936607, ...,  1.4842858 ,
      -0.01975653,  0.36723647],
       [ 0.08518241, -0.94217861,  1.01898711, ..., -0.06221033,
       0.3312366 , -0.31866022],
       ...,
       [ 0.63475351, -0.94217861, -0.46502808, ..., -0.83545839,
      -0.25375196, -0.06144896],
       [-0.30736838, -0.94217861, -0.53352109, ..., -0.06221033,
      -0.83072436, -0.83308273],
       [-2.03459183, -0.94217861,  0.56236706, ..., -0.83545839,
      -0.13312789, -0.06144896]], shape=(353, 10))
```


```
[6]: y_pred=model.predict(x_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```

```
[6]: 0.7752808988764045
```

```
[7]: print(confusion_matrix(y_test,y_pred))
```

```
[[35 14]
 [ 6 34]]
```

```
[8]: # Visualize the decision boundary with accuracy information
plt.figure(figsize=(8, 6))
sns.scatterplot(x=x_test[:, 2], y=x_test[:, 8], hue=y_test, palette={
    0: 'blue', 1: 'red'}, marker='o')
plt.xlabel("BMI")
plt.ylabel("Age")
plt.title("Logistic Regression Decision Boundary\nAccuracy: {:.2f}%".format(
    accuracy * 100))
plt.legend(title="Diabetes", loc="upper right")
plt.show()
```



## PROGRAM 4 : (DECISION TREE)

```
[6]: import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier,plot_tree
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
import matplotlib.pyplot as plt

[7]: df=pd.read_csv("adult.csv")

[8]: df.replace('?',np.nan,inplace=True)
df.dropna(inplace=True)
df.drop(['fnlwgt','educational-num','marital-status','relationship','race'],axis=1,inplace=True)

[9]: X=df.drop('income',axis=1)
y=LabelEncoder().fit_transform(df['income'])

[10]: X=pd.get_dummies(X)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.33,random_state=42)
clf=DecisionTreeClassifier(max_depth=5)
clf.fit(X_train,y_train)
y_pred=clf.predict(X_test)
print("Model Accuracy:",accuracy_score(y_test,y_pred))


Model Accuracy: 0.8234387563655856

[11]: cm=confusion_matrix(y_test,y_pred)
print("\n Confusion Matrix:\n",cm)
print("\n Classification Report:\n",classification_report(y_test,y_pred))

Confusion Matrix:
[[10736  496]
 [ 2139 1553]]
```

```
[12]: plt.figure(figsize=(14,14))
plot_tree(clf,fontsize=10,filled=True)
plt.title("Decision Tree trained on adult income dataset")
plt.show()
```

Decision Tree trained on adult income dataset



## PROGRAM 5 : (SVM)

```
[4]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import LabelEncoder
```

```
[5]: df=pd.read_csv("smaller_adult.csv")
```

```
[6]: df.replace("?",np.nan,inplace=True)
df.dropna(inplace=True)
df.drop_duplicates(inplace=True)
X=df[['age','workclass','educational-num','occupation','gender','hours-per-week']]
y=LabelEncoder().fit_transform(df['income'])
```

```
[7]: X=pd.get_dummies(X)
```

```
[8]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.33,random_state=42)
```

```
[9]: clf=SVC(kernel='linear',gamma='auto')
clf.fit(X_train,y_train)
```

```
[9]: ▾ SVC
SVC(gamma='auto', kernel='linear')
```

```
[10]: y_pred=clf.predict(X_test)
print("Accuracy:",accuracy_score(y_test,y_pred))
```

Accuracy: 0.7672162948593598